

SEMANTIC FORCE CLOSURE REVISITED: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study grasping and language-conditioned manipulation. The motivating bet is: Test when semantic object categories change the relevant closure conditions.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For grasping and language-conditioned manipulation, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2023a; 2026a;b;c; 2023b;c). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.562	0.438
observed_only_wam	hidden_mode_shift	0.147	0.853
augmentation_as_data_wam	nominal	0.593	0.407
augmentation_as_data_wam	hidden_mode_shift	0.228	0.772
scalar_uncertainty_wam	nominal	0.675	0.325
scalar_uncertainty_wam	hidden_mode_shift	0.280	0.720
proposed_branch_mechanism	nominal	0.782	0.218
proposed_branch_mechanism	hidden_mode_shift	0.588	0.412

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Semantic Force Closure Revisited argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Pong: Probabilistic object normals for grasping via analytic bounds on force closure probability, 2023a. <http://arxiv.org/abs/2309.16930v1>.
- Analysis of contact force and shape change on grasping a square object using an actual fin ray soft gripper, 2023b. <https://doi.org/10.3390/s23249827>.
- Lan-grasp: Using large language models for semantic object grasping and placement, 2023c. <http://arxiv.org/abs/2310.05239v3>.
- Graspmind: Multi-granularity language-conditioned affordance prediction for robotic grasping, 2026a. <https://doi.org/10.2139/ssrn.6365601>.
- Development and evaluation of a vision-language guided robotic manipulation pipeline, 2026b. <https://doi.org/10.32920/32253735.v1>.
- Development and evaluation of a vision-language guided robotic manipulation pipeline, 2026c. <https://doi.org/10.32920/32253735>.